

# Backend Engineer - Coding Tasks

## Abstract

The task is intended to expose the candidate to small-scope real-life design and development challenges.

## Assumptions

1. You can use any backend Java framework you like e.g. Spring Boot, J2EE, J2EE Microprofile, OSGI (god forbid ;))
2. If you need a DB, please use (for the task) something embedded like H2, HSQLDB, or Derby.
3. You can (and should, if possible) use third-party libraries to leverage parts of the functionality that are “generic.”

## Domain Background

## Vocabulary

A **Treatment Plan** describes the type and schedule of the treatment actions for a particular patient and has the following properties:

1. Treatment action: enum of possible actions [ActionA or ActionB]
2. Subject patient: reference to the subject patient
3. Start time: timestamp of earliest possible treatment task
4. End time: timestamp of the latest possible treatment task or empty if the plan is endless
5. Recurrence pattern: for example, “every day at 08:00 and 16:00” or another example, “every Monday at 12:00”

A **Treatment Task** describes a single treatment action to be done for a specific patient at a specific time and has the following properties:

1. Treatment action: enum of possible actions [ActionA or ActionB]
2. Subject patient: reference to the subject patient
3. Start time: timestamp of the earliest time the task can be completed
4. Status: Active or Completed

# Functional Requirements

1. [Given, outside of scope, **no need to implement**] As a Clinician, I can create a treatment plan for a specific patient so I can communicate to all the treatment tasks executor what has to be done
2. [Given, outside of scope, **no need to implement**] As a Nurse, I can browse/view treatment tasks so I can execute treatment actions on time and report their completion
3. **[Need to be implemented]** As a System, I can generate treatment tasks based on the treatment plan
  - a. Your goal is to implement a “scheduling” execution service that “spawns” treatment tasks according to all the treatment plans existing in the system and their schedule (start time, recurrence pattern, and end time)
  - b. We can **start with simple implementation** and then discuss the non-functional requirements like scalability and resilience.
  - c. Software Design Hints
    - i. Treatment Plans and Treatment Tasks are stored in separate tables in DB, but you decide on the schema/layout as needed.
      1. Reference to the patient is a simple string
      2. The entities are denormalized, and this is OK
      3. You can add any field you think is needed.
    - ii. Treatment Plan’s insertion into DB is outside of the scope of this task, and you can assume that they are already there
    - iii. The result of this coding task is a scheduler that creates new Treatment Tasks stored in DB according to what is defined in active Treatment Plans it can find in DB. No need to execute the tasks; just store them in DB.
    - iv. You should not use the Quartz scheduler as a whole; you are supposed to design and implement the “executor” part of its functionality.

## FAQ and Comments:

1. Should I support human-readable recurrence patterns? - Recurrence patterns format is an implementation detail and is not required to be humanly readable
2. Should I create all or some of the treatment tasks ahead of time? - No need. You need the treatment task to exist in the DB when the task is active. Creating too many tasks ahead of time can be a scalability issue.